

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
from sklearn.cluster import KMeans
import warnings
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import StandardScaler, PolynomialFeatures

# Ignore FutureWarnings
warnings.simplefilter(action='ignore', category=FutureWarning)

# Load dataset
houseprice_df = pd.read_excel('HousePrice_Dataset.xlsx')
houseprice_df.head()
```

```
Out[1]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-122.23	37.88	41	880	129.0	322	126	8.3
1	-122.22	37.86	21	7099	1106.0	2401	1138	8.3
2	-122.25	37.85	52	1627	280.0	565	259	3.8
3	-122.25	37.85	52	919	213.0	413	193	4.0
4	-122.25	37.84	52	2535	489.0	1094	514	3.6

```
In [2]: # Checking input data shape & info
df = houseprice_df.copy()
print(df.shape)
print(df.info())

(18565, 10)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18565 entries, 0 to 18564
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              18565 non-null  float64
1   latitude               18565 non-null  float64
2   housing_median_age     18565 non-null  int64
3   total_rooms            18565 non-null  int64
4   total_bedrooms         18376 non-null  float64
5   population             18565 non-null  int64
6   households             18565 non-null  int64
7   median_income          18565 non-null  float64
8   median_house_value     18565 non-null  int64
9   ocean_proximity        18565 non-null  object
dtypes: float64(4), int64(5), object(1)
memory usage: 1.4+ MB
None
```

```
In [3]: #Checking for NULL values in Data
df.isnull().sum()
```

```
Out[3]: longitude      0
latitude      0
housing_median_age  0
total_rooms    0
total_bedrooms 189
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64
```

```
In [4]: #Dropping NULL Values
df = df.dropna()
```

```
In [5]: #Confirming dropna operation
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 18376 entries, 0 to 18564
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             18376 non-null  float64
1   latitude              18376 non-null  float64
2   housing_median_age    18376 non-null  int64
3   total_rooms           18376 non-null  int64
4   total_bedrooms        18376 non-null  float64
5   population            18376 non-null  int64
6   households            18376 non-null  int64
7   median_income         18376 non-null  float64
8   median_house_value    18376 non-null  int64
9   ocean_proximity       18376 non-null  object
dtypes: float64(4), int64(5), object(1)
memory usage: 1.5+ MB
```

```
In [6]: #Checking for any duplicates
df.duplicated().sum()
```

```
Out[6]: 0
```

```
In [7]: #Checking data composition for numerical features
df.describe()
```

```
Out[7]:
```

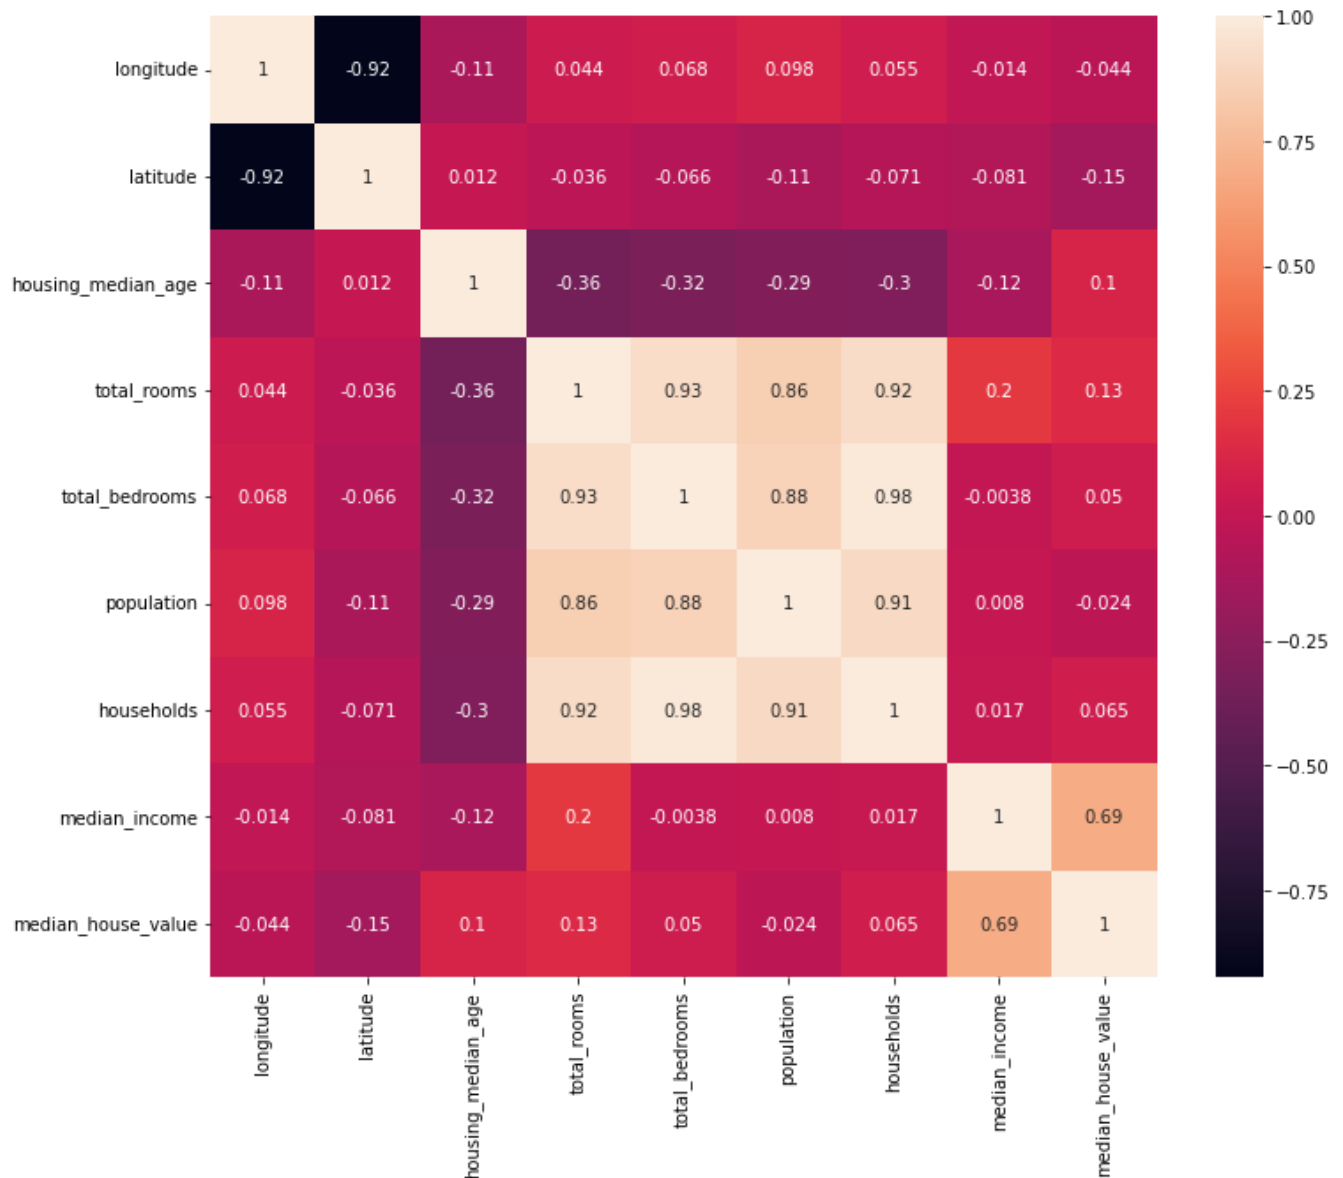
	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	househ
count	18376.000000	18376.000000	18376.000000	18376.000000	18376.000000	18376.000000	18376.00
mean	-119.571095	35.635164	28.605736	2635.302188	537.711199	1425.810786	499.37
std	2.003042	2.137485	12.570789	2200.534974	424.125849	1143.481721	384.51
min	-124.350000	32.540000	1.000000	2.000000	2.000000	3.000000	2.00
25%	-121.800000	33.930000	18.000000	1444.000000	295.000000	786.000000	280.00
50%	-118.500000	34.260000	29.000000	2123.000000	434.000000	1165.500000	408.00
75%	-118.010000	37.720000	37.000000	3137.000000	646.000000	1722.000000	603.00
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.00

```
In [8]: #Standardizing the numerical feilds.
# df[['housing_median_age']] = preprocessing.scale(df[['housing_median_age']]).astype('float64')
# df[['total_rooms']] = preprocessing.scale(df[['total_rooms']]).astype('float64')
```

```
# df[['total_bedrooms']] = preprocessing.scale(df[['total_bedrooms']]).astype('float64')
# df[['population']] = preprocessing.scale(df[['population']]).astype('float64')
# df[['households']] = preprocessing.scale(df[['households']]).astype('float64')
# df[['median_income']] = preprocessing.scale(df[['median_income']]).astype('float64')
# df[['median_house_value']] = preprocessing.scale(df[['median_house_value']]).astype('float64')


# df.head()
```

```
In [9]: # Correlation matrix
figure1, a = plt.subplots(figsize=(12, 10))
sns.heatmap(df.corr(), annot=True)
plt.show()
```



```
In [10]: df=df.drop(['total_rooms','total_bedrooms','population'], axis=1)
```

```
In [11]: # Correlation matrix
figure1, a = plt.subplots(figsize=(12, 10))
sns.heatmap(df.corr(), annot=True)
plt.show()
```



```
In [12]: #Checking unique values in categorical column
df['ocean_proximity'].unique()
```


```
Out[12]: array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
      dtype=object)
```

```
In [13]: #House price based on categorical value.
df.groupby(['ocean_proximity'])['median_house_value'].describe()
```

```
Out[13]:
```


	count	mean	std	min	25%	50%	75%	max
<b>ocean_proximity</b>								
<1H OCEAN	8096.0	239973.465415	106101.457273	17500.0	164175.0	214750.0	289425.0	500001.0
INLAND	5869.0	124937.335492	70783.931014	14999.0	77500.0	108300.0	148600.0	500001.0
ISLAND	5.0	380440.000000	80559.561816	287500.0	300000.0	414700.0	450000.0	450000.0
NEAR BAY	2034.0	258756.622911	122646.084078	22500.0	162500.0	231800.0	345525.0	500001.0
NEAR OCEAN	2372.0	249858.342327	122701.540906	22500.0	150000.0	229800.0	323825.0	500001.0

```
In [14]: #Same information using a bar chart.
df['ocean_proximity'].value_counts().plot(kind='bar')
plt.show()
```




```
In [15]: figure2, ax = plt.subplots(figsize = (12,8))
sns.barplot(x = 'ocean_proximity', y = 'median_house_value' , data = df)
```

```
Out[15]: <AxesSubplot:xlabel='ocean_proximity', ylabel='median_house_value'>
```



```
In [16]: figure3, ax = plt.subplots(figsize = (12,8))
sns.scatterplot(x = 'median_income', y = 'median_house_value' , data = df)
```

```
Out[16]: <AxesSubplot:xlabel='median_income', ylabel='median_house_value'>
```



```
In [17]: # Convert categorical variable 'ocean_proximity' into dummy/indicator variables
df = pd.get_dummies(df, columns=['ocean_proximity'])
```


```
In [18]: #Renaming feature name as it should not contain characters like <=> etc...This will give problem
df.rename(columns = {'ocean_proximity_<1H OCEAN' : 'ocean_proximity_LT1H OCEAN'}, inplace = True)
```

```
In [19]: #Checking for dummies action & feature name change
df.head()
```

```
Out[19]:
```


	longitude	latitude	housing_median_age	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41	126	8.3252	452600	INLAND
1	-122.22	37.86	21	1138	8.3014	358500	INLAND
2	-122.25	37.85	52	259	3.8462	342200	INLAND
3	-122.25	37.85	52	193	4.0368	269700	INLAND
4	-122.25	37.84	52	514	3.6591	299200	INLAND

```
In [20]: #Distribution of median house value which we are predicting
#Data shows right skewed and some outliers
sns.distplot(df['median_house_value'])
plt.show()
```



```
In [21]: #Confirming via box plot
sns.boxplot(df['median_house_value'])
```

```
Out[21]: <AxesSubplot:xlabel='median_house_value'>
```



```
In [22]: #Trying to identify outliers and remove them
# IQR
# Calculate the upper and lower limits
Q1 = df['median_house_value'].quantile(0.25)
Q3 = df['median_house_value'].quantile(0.75)
IQR = Q3 - Q1

# Set the threshold for outliers
lower_threshold = Q1 - 1.5 * IQR
upper_threshold = Q3 + 1.5 * IQR

# Identify outlier rows using boolean masks
outliers_lower = df['median_house_value'] < lower_threshold
outliers_upper = df['median_house_value'] > upper_threshold


# Display the number of outliers
print(f"Number of lower outliers: {outliers_lower.sum()}")
print(f"Number of upper outliers: {outliers_upper.sum()}")

# Remove the outliers
df = df[~(outliers_lower | outliers_upper)].copy()

# Display the shape of the cleaned DataFrame
print(f"Shape of the cleaned DataFrame: {df.shape}")
```


Number of lower outliers: 0  
Number of upper outliers: 957  
Shape of the cleaned DataFrame: (17419, 11)

```
In [23]: #Confirming outliers removal
sns.distplot(df['median_house_value'])
plt.show()
```



```
In [24]: sns.distplot(np.log(df['median_house_value']))
```

```
Out[24]: <AxesSubplot:xlabel='median_house_value', ylabel='Density'>
```



```
In [25]: #Checking if any column belongs to object
category_columns = df.dtypes[df.dtypes=='object'].index
numerical_columns = df.dtypes[df.dtypes!='object'].index
print(category_columns)
print(numerical_columns)
```

```
Index([], dtype='object')
Index(['longitude', 'latitude', 'housing_median_age', 'households',
       'median_income', 'median_house_value', 'ocean_proximity_LT1H OCEAN',
       'ocean_proximity_INLAND', 'ocean_proximity_ISLAND',
       'ocean_proximity_NEAR BAY', 'ocean_proximity_NEAR OCEAN'],
      dtype='object')
```

```
In [26]: #Seperating input features and Labels
X = df.drop('median_house_value', axis = 1)
Y = np.log(df['median_house_value'])
print(type(X))
print(type(Y))
```



```
print(X.shape)
print(Y.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
(17419, 10)
(17419,)
```

```
In [27]: # Split the dataset into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
df.head()
```

```
(13935, 10)
(3484, 10)
(13935,)
(3484,)
```

```
Out[27]:
```

	longitude	latitude	housing_median_age	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41	126	8.3252	452600	C
1	-122.22	37.86	21	1138	8.3014	358500	
2	-122.25	37.85	52	259	3.8462	342200	
3	-122.25	37.85	52	193	4.0368	269700	
4	-122.25	37.84	52	514	3.6591	299200	

```
In [28]: ## Standardize the features using StandardScaler
# scaler = StandardScaler()
# X_train = scaler.fit_transform(x_train)
# X_test = scaler.transform(x_test)
```

```
In [29]: #Defining functions to calculate R2 and error rate
def model_evaluation(y_test, y_pred):
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = mean_squared_error(y_test, y_pred, squared=False)
    # r2_scr = r2_score(y_test, y_pred)
    return {'mae': mae, 'mse': mse, 'rmse': rmse}

def model_instantiation(model, x_train, x_test, y_train, y_test, y_pred, model_name):
    training_r2 = model.score(x_train, y_train)
    testing_r2 = model.score(x_test, y_test)
    eval_return = model_evaluation(y_test, y_pred)
    result_metric = {
        'Train_R2': training_r2,
        'Test_R2': testing_r2,
        'Test_MSE': eval_return['mse'],
        'Test_RMSE': eval_return['rmse'],
        'Test_MAE': eval_return['mae']
    }
    result = pd.DataFrame(result_metric, index=[model_name])
    return result
```

```
In [30]: from sklearn.linear_model import LinearRegression
```

```
In [31]: #Instantiating Linear reg model
linear_reg = LinearRegression()
```

```
linear_reg.fit(x_train, y_train)
```

```
Out[31]: ▼ LinearRegression  
LinearRegression()
```

```
In [32]: #Predicting via LR  
y_pred_LR = linear_reg.predict(x_test)
```

```
In [33]: LR_df = model_instantiation(linear_reg ,x_train,x_test,y_train,y_test,y_pred_LR , 'Linear Reg  
LR_df
```

```
Out[33]:
```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>Linear Regression</b>	0.630962	0.624499	0.10712	0.327293	0.251028

```
In [34]: from sklearn.tree import DecisionTreeRegressor  
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor
```

```
In [35]: #Decision Tree  
Decs_Tree = DecisionTreeRegressor(max_depth = 8 , min_samples_leaf =10 , min_samples_split =  
Decs_Tree.fit(x_train, y_train)
```

```
Out[35]: ▼ DecisionTreeRegressor  
DecisionTreeRegressor(max_depth=8, min_samples_leaf=10, min_samples_split=10)
```

```
In [36]: y_pred_DT = Decs_Tree.predict(x_test)
```

```
In [37]: DT_df = model_instantiation(Decs_Tree ,x_train,x_test,y_train,y_test,y_pred_DT , 'DTree_Regre  
DT_df
```

```
Out[37]:
```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>DTree_Regression</b>	0.763267	0.717443	0.080606	0.283912	0.208418

```
In [38]: #Random forest  
Rand_Forest = RandomForestRegressor(n_estimators=300, max_depth = 10, min_samples_split=12)  
Rand_Forest.fit(x_train, y_train)
```

```
Out[38]: ▼ RandomForestRegressor  
RandomForestRegressor(max_depth=10, min_samples_split=12, n_estimators=300)
```

```
In [39]: y_pred_RF = Rand_Forest.predict(x_test)
```

```
In [40]: RF_df = model_instantiation(Rand_Forest ,x_train,x_test,y_train,y_test,y_pred_RF , 'RF_Regres  
RF_df
```

```
Out[40]:
```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>RF_Regression</b>	0.850595	0.774959	0.064198	0.253374	0.180186

```
In [41]: Rand_Forest_2 = RandomForestRegressor(n_estimators=300, max_depth = 9, min_samples_split=4)  
Rand_Forest_2.fit(x_train, y_train)  
  
y_pred_RF_2 = Rand_Forest_2.predict(x_test)
```

```
RF_df_2 = model_instantiation(Rand_Forest_2 ,x_train,x_test,y_train,y_test,y_pred_RF_2 , 'RF2')
RF_df_2
```

```
Out[41]:
```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>RF2_Regression</b>	0.830124	0.763076	0.067588	0.259977	0.186851

```
In [42]: ada =AdaBoostRegressor(n_estimators = 300, random_state = 10)
ada.fit(x_train,y_train)

y_pred_ada = ada.predict(x_test)

ada_df = model_instantiation(ada ,x_train,x_test,y_train,y_test,y_pred_ada , 'ADA_Regression')
ada_df
```

```
Out[42]:
```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>ADA_Regression</b>	0.599874	0.586475	0.117968	0.343464	0.271343

```
In [43]: from xgboost.sklearn import XGBRegressor

xgb_model = XGBRegressor()

# Define the parameter grid to search
# param_grid = {
#     'n_estimators': [100, 200, 300],
#     'learning_rate': [0.01, 0.1, 0.2],
#     'max_depth': [3, 4, 5],
#     'subsample': [0.8, 0.9, 1.0],
#     'colsample_bytree': [0.8, 0.9, 1.0],
#     'reg_alpha': [0, 0.1, 0.5], # Regularization term on weights (L1)
#     'reg_lambda': [0, 0.1, 0.5] # Regularization term on weights (L2)
# }

# # # Create GridSearchCV
# grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, scoring='neg_mean_sq

# grid_search.fit(x_train,y_train,early_stopping_rounds=10,eval_set=[(x_test, y_test)], verbo
# XGB_Best_Model = grid_search.best_estimator_
# y_pred_xgboost = XGB_Best_Model.predict(x_test)

xgb_model.fit(x_train,y_train)
y_pred_xgboost = xgb_model.predict(x_test)

xgboost_df = model_instantiation(xgb_model, x_train, x_test, y_train, y_test, y_pred_xgboost,
xgboost_df
```

```
Out[43]:
```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>XGB_Regression</b>	0.923149	0.801669	0.056579	0.237862	0.163638

```
In [ ]:
```

```
In [44]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.pipeline import make_pipeline
from tensorflow.keras.regularizers import l2

# # Separate features(X) and target variable (y)
X = df.drop('median_house_value', axis=1)
```

```

y = df['median_house_value']

## Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

## Create a pipeline with StandardScaler and PolynomialFeatures
pipeline = make_pipeline(StandardScaler(), PolynomialFeatures(degree=2, include_bias=False))

## Fit and transform the training data
X_train_poly = pipeline.fit_transform(X_train)

## Transform the testing data
X_test_poly = pipeline.transform(X_test)

# Build a multi-layer neural network model
model = Sequential()
model.add(Dense(128, input_dim=X_train_poly.shape[1], activation='relu', kernel_regularizer=l2(0.001)))
model.add(Dropout(0.2))
model.add(Dense(64, activation='relu', kernel_regularizer=l2(0.001)))
model.add(BatchNormalization())
model.add(Dense(32, activation='relu', kernel_regularizer=l2(0.001)))
model.add(Dense(16, activation='relu', kernel_regularizer=l2(0.001)))
model.add(Dense(1, activation='relu'))

# Define early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)

# Compile the model with an adjusted learning rate
optimizer = tf.keras.optimizers.Adam(learning_rate=0.0001)
model.compile(optimizer=optimizer, loss='mean_squared_error')

# Train the model
model.fit(X_train_poly, y_train, epochs=50, batch_size=32, validation_split=0.2, callbacks=[early_stopping])

# Make predictions on the test set
predictions = model.predict(X_test_poly)

# Evaluate the model
mse = mean_squared_error(y_test, predictions)
r2 = r2_score(y_test, predictions)

print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')
# xgboost_df = model_instantiation(model, x_train, x_test, y_train, y_test, predictions, 'Ten')
# xgboost_df

```

Epoch 1/50  
349/349 [=====] - 3s 4ms/step - loss: 45521571840.0000 - val\_loss: 4  
5659099136.0000

Epoch 2/50  
349/349 [=====] - 1s 4ms/step - loss: 45513981952.0000 - val\_loss: 4  
5651877888.0000

Epoch 3/50  
349/349 [=====] - 1s 3ms/step - loss: 45483630592.0000 - val\_loss: 4  
5618933760.0000

Epoch 4/50  
349/349 [=====] - 1s 3ms/step - loss: 45420085248.0000 - val\_loss: 4  
5588664320.0000

Epoch 5/50  
349/349 [=====] - 1s 3ms/step - loss: 45316804608.0000 - val\_loss: 4  
5505753088.0000

Epoch 6/50  
349/349 [=====] - 1s 3ms/step - loss: 45167734784.0000 - val\_loss: 4  
5382037504.0000

Epoch 7/50  
349/349 [=====] - 1s 3ms/step - loss: 44966608896.0000 - val\_loss: 4  
5313585152.0000

Epoch 8/50  
349/349 [=====] - 1s 3ms/step - loss: 44703805440.0000 - val\_loss: 4  
5178540032.0000

Epoch 9/50  
349/349 [=====] - 1s 3ms/step - loss: 44380065792.0000 - val\_loss: 4  
5011771392.0000

Epoch 10/50  
349/349 [=====] - 1s 3ms/step - loss: 43981303808.0000 - val\_loss: 4  
4739723264.0000

Epoch 11/50  
349/349 [=====] - 1s 3ms/step - loss: 43510579200.0000 - val\_loss: 4  
4404719616.0000

Epoch 12/50  
349/349 [=====] - 1s 2ms/step - loss: 42954420224.0000 - val\_loss: 4  
4059979776.0000

Epoch 13/50  
349/349 [=====] - 1s 3ms/step - loss: 42316414976.0000 - val\_loss: 4  
3515187200.0000

Epoch 14/50  
349/349 [=====] - 1s 3ms/step - loss: 41571135488.0000 - val\_loss: 4  
3261763584.0000

Epoch 15/50  
349/349 [=====] - 1s 3ms/step - loss: 40685912064.0000 - val\_loss: 4  
2615050240.0000

Epoch 16/50  
349/349 [=====] - 1s 3ms/step - loss: 39722323968.0000 - val\_loss: 4  
1998520320.0000

Epoch 17/50  
349/349 [=====] - 1s 3ms/step - loss: 38568656896.0000 - val\_loss: 4  
1106513920.0000

Epoch 18/50  
349/349 [=====] - 1s 3ms/step - loss: 37363703808.0000 - val\_loss: 4  
0114827264.0000

Epoch 19/50  
349/349 [=====] - 1s 4ms/step - loss: 36029636608.0000 - val\_loss: 3  
9349116928.0000

Epoch 20/50  
349/349 [=====] - 1s 3ms/step - loss: 34568691712.0000 - val\_loss: 3  
8629601280.0000

Epoch 21/50  
349/349 [=====] - 1s 3ms/step - loss: 33029892096.0000 - val\_loss: 3  
8276018176.0000

Epoch 22/50  
349/349 [=====] - 1s 3ms/step - loss: 31390441472.0000 - val\_loss: 3  
6098195456.0000

Epoch 23/50  
349/349 [=====] - 1s 3ms/step - loss: 29657585664.0000 - val\_loss: 3  
5949268992.0000

Epoch 24/50  
349/349 [=====] - 1s 3ms/step - loss: 27854807040.0000 - val\_loss: 3  
4343340032.0000

Epoch 25/50  
349/349 [=====] - 1s 3ms/step - loss: 26037495808.0000 - val\_loss: 3  
2776638464.0000

Epoch 26/50  
349/349 [=====] - 1s 3ms/step - loss: 24125790208.0000 - val\_loss: 3  
1230275584.0000

Epoch 27/50  
349/349 [=====] - 1s 3ms/step - loss: 22231838720.0000 - val\_loss: 3  
0732843008.0000

Epoch 28/50  
349/349 [=====] - 1s 2ms/step - loss: 20321949696.0000 - val\_loss: 2  
8998916096.0000

Epoch 29/50  
349/349 [=====] - 1s 3ms/step - loss: 18474164224.0000 - val\_loss: 2  
7636432896.0000

Epoch 30/50  
349/349 [=====] - 1s 4ms/step - loss: 16617832448.0000 - val\_loss: 2  
5972879360.0000

Epoch 31/50  
349/349 [=====] - 1s 3ms/step - loss: 15004934144.0000 - val\_loss: 2  
5463404544.0000

Epoch 32/50  
349/349 [=====] - 1s 4ms/step - loss: 13281187840.0000 - val\_loss: 2  
4184819712.0000

Epoch 33/50  
349/349 [=====] - 1s 3ms/step - loss: 12014547968.0000 - val\_loss: 2  
1572216832.0000

Epoch 34/50  
349/349 [=====] - 1s 3ms/step - loss: 10627654656.0000 - val\_loss: 2  
2135050240.0000

Epoch 35/50  
349/349 [=====] - 1s 3ms/step - loss: 9489682432.0000 - val\_loss: 20  
834973696.0000

Epoch 36/50  
349/349 [=====] - 1s 4ms/step - loss: 8421891584.0000 - val\_loss: 19  
097423872.0000

Epoch 37/50  
349/349 [=====] - 1s 4ms/step - loss: 7558081024.0000 - val\_loss: 19  
530541056.0000

Epoch 38/50  
349/349 [=====] - 1s 3ms/step - loss: 6951711744.0000 - val\_loss: 18  
991861760.0000

Epoch 39/50  
349/349 [=====] - 1s 3ms/step - loss: 6637931008.0000 - val\_loss: 16  
955921408.0000

Epoch 40/50  
349/349 [=====] - 1s 3ms/step - loss: 6067017216.0000 - val\_loss: 15  
321756672.0000

Epoch 41/50  
349/349 [=====] - 1s 3ms/step - loss: 5854266880.0000 - val\_loss: 14  
848336896.0000

Epoch 42/50  
349/349 [=====] - 1s 3ms/step - loss: 5741652992.0000 - val\_loss: 13  
548585984.0000

Epoch 43/50  
349/349 [=====] - 1s 3ms/step - loss: 5499480064.0000 - val\_loss: 13  
736945664.0000

Epoch 44/50  
349/349 [=====] - 1s 3ms/step - loss: 5346961408.0000 - val\_loss: 13  
284270080.0000

```

Epoch 45/50
349/349 [=====] - 1s 4ms/step - loss: 5179998720.0000 - val_loss: 12
758976512.0000
Epoch 46/50
349/349 [=====] - 1s 3ms/step - loss: 5294595072.0000 - val_loss: 12
697150464.0000
Epoch 47/50
349/349 [=====] - 1s 3ms/step - loss: 5125353472.0000 - val_loss: 11
694658560.0000
Epoch 48/50
349/349 [=====] - 1s 3ms/step - loss: 5200043520.0000 - val_loss: 97
71138048.0000
Epoch 49/50
349/349 [=====] - 1s 3ms/step - loss: 5213671936.0000 - val_loss: 94
03393024.0000
Epoch 50/50
349/349 [=====] - 1s 3ms/step - loss: 5141705728.0000 - val_loss: 10
682987520.0000
109/109 [=====] - 1s 1ms/step
Mean Squared Error: 10499794813.146599
R-squared: -0.1518082573703612

```

In [45]: `from sklearn.linear_model import Ridge, Lasso`

```

lasso = Lasso(alpha=1.0)
lasso.fit(x_train, y_train)
y_pred_lasso = lasso.predict(x_test)

```

```

lasso_df = model_instantiation(lasso, x_train, x_test, y_train, y_test, y_pred_lasso, 'Lasso')
lasso_df

```

```

C:\Users\sumit\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:628: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 4.345e+12, tolerance: 1.269e+10
model = cd_fast.enet_coordinate_descent(

```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>Lasso</b>	0.585009	0.588948	3.747114e+09	61213.678646	46199.161283

In [46]: `ridge = Ridge(alpha=1.0)`  
`ridge.fit(x_train, y_train)`  
`y_pred_ridge = ridge.predict(x_test)`

```

ridge_df = model_instantiation(ridge, x_train, x_test, y_train, y_test, y_pred_ridge, 'Ridge')
ridge_df

```

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
<b>Ridge</b>	0.584982	0.588941	3.747182e+09	61214.231846	46202.274281

In [47]: `all_results = pd.concat([LR_df, DT_df, RF_df, RF_df_2, ada_df, xgboost_df, lasso_df, ridge_df])`  
`all_results`

Out[47]:

	Train_R2	Test_R2	Test_MSE	Test_RMSE	Test_MAE
Linear Regression	0.630962	0.624499	1.071205e-01	0.327293	0.251028
DTree_Regression	0.763267	0.717443	8.060598e-02	0.283912	0.208418
RF_Regression	0.850595	0.774959	6.419829e-02	0.253374	0.180186
RF2_Regression	0.830124	0.763076	6.758802e-02	0.259977	0.186851
ADA_Regression	0.599874	0.586475	1.179678e-01	0.343464	0.271343
XGB_Regression	0.923149	0.801669	5.657852e-02	0.237862	0.163638
Lasso	0.585009	0.588948	3.747114e+09	61213.678646	46199.161283
Ridge	0.584982	0.588941	3.747182e+09	61214.231846	46202.274281